

# Complete Machine Learning Cheatsheet

## Table of Contents

1. Types of Machine Learning Problems
  2. Supervised Learning Algorithms
  3. Unsupervised Learning Algorithms
  4. Reinforcement Learning
  5. Data Preprocessing Techniques
  6. Feature Engineering
  7. Model Evaluation Metrics
  8. Model Selection & Validation
  9. Common Challenges & Solutions
  10. Deep Learning Overview
- 

## Types of Machine Learning Problems

### 1. Supervised Learning

**Definition:** Learning from labeled data to predict outcomes for new data.

#### Classification

- **Binary Classification:** Two classes (e.g., spam/not spam)
- **Multi-class Classification:** Multiple classes (e.g., digit recognition 0-9)
- **Multi-label Classification:** Multiple labels per instance

#### Regression

- **Linear Regression:** Predicting continuous values
- **Non-linear Regression:** Complex relationships between variables

### 2. Unsupervised Learning

**Definition:** Finding patterns in unlabeled data.

- **Clustering:** Grouping similar data points
- **Dimensionality Reduction:** Reducing feature space
- **Anomaly Detection:** Identifying outliers
- **Association Rules:** Finding relationships between variables

### 3. Semi-Supervised Learning

**Definition:** Combination of labeled and unlabeled data.

### 4. Reinforcement Learning

**Definition:** Learning through interaction with environment to maximize reward.

---

## Supervised Learning Algorithms

### Classification Algorithms

#### 1. Logistic Regression

- **Use Case:** Binary/multi-class classification
- **Pros:** Simple, interpretable, probabilistic output
- **Cons:** Assumes linear relationship
- **Key Parameters:** C (regularization), solver, penalty

#### 2. Decision Trees

- **Use Case:** Both classification and regression
- **Pros:** Interpretable, handles non-linear relationships
- **Cons:** Prone to overfitting
- **Key Parameters:** max\_depth, min\_samples\_split, criterion

#### 3. Random Forest

- **Use Case:** Complex classification/regression problems
- **Pros:** Reduces overfitting, handles missing values
- **Cons:** Less interpretable, computationally expensive
- **Key Parameters:** n\_estimators, max\_depth, min\_samples\_split

#### 4. Support Vector Machines (SVM)

- **Use Case:** High-dimensional data, non-linear problems
- **Pros:** Effective in high dimensions, memory efficient
- **Cons:** Slow on large datasets, requires feature scaling
- **Key Parameters:** C, kernel, gamma

#### 5. k-Nearest Neighbors (k-NN)

- **Use Case:** Simple classification/regression

- **Pros:** Simple, no training phase
- **Cons:** Slow prediction, sensitive to feature scale
- **Key Parameters:** n\_neighbors, weights, metric

## 6. Naive Bayes

- **Use Case:** Text classification, spam filtering
- **Pros:** Fast, works well with high dimensions
- **Cons:** Assumes feature independence
- **Types:** Gaussian, Multinomial, Bernoulli

## 7. Gradient Boosting (XGBoost, LightGBM, CatBoost)

- **Use Case:** Competitions, high-performance needs
- **Pros:** High accuracy, handles various data types
- **Cons:** Prone to overfitting, many hyperparameters
- **Key Parameters:** learning\_rate, n\_estimators, max\_depth

## Regression Algorithms

### 1. Linear Regression

- **Use Case:** Linear relationships
- **Pros:** Simple, interpretable
- **Cons:** Assumes linearity
- **Variants:** Ridge, Lasso, Elastic Net

### 2. Polynomial Regression

- **Use Case:** Non-linear relationships
- **Pros:** Flexible
- **Cons:** Prone to overfitting

### 3. Support Vector Regression (SVR)

- **Use Case:** Non-linear regression
- **Pros:** Robust to outliers
- **Cons:** Computationally expensive

---

## Unsupervised Learning Algorithms

### Clustering Algorithms

## 1. K-Means

- **Use Case:** Well-separated spherical clusters
- **Pros:** Fast, scalable
- **Cons:** Requires k specification, sensitive to initialization
- **Key Parameters:** n\_clusters, init, max\_iter

## 2. Hierarchical Clustering

- **Use Case:** Hierarchical data structure
- **Pros:** No need to specify k, dendrogram visualization
- **Cons:** Computationally expensive
- **Types:** Agglomerative, Divisive

## 3. DBSCAN

- **Use Case:** Arbitrary shaped clusters, noise handling
- **Pros:** No need to specify k, finds outliers
- **Cons:** Sensitive to parameters
- **Key Parameters:** eps, min\_samples

## 4. Gaussian Mixture Models (GMM)

- **Use Case:** Soft clustering, overlapping clusters
- **Pros:** Probabilistic approach
- **Cons:** Assumes Gaussian distribution

## Dimensionality Reduction

### 1. Principal Component Analysis (PCA)

- **Use Case:** Linear dimensionality reduction
- **Pros:** Fast, well-understood
- **Cons:** Linear only, loses interpretability

### 2. t-SNE

- **Use Case:** Visualization of high-dimensional data
- **Pros:** Excellent for visualization
- **Cons:** Computationally expensive, not for new data

### 3. UMAP

- **Use Case:** General dimensionality reduction
- **Pros:** Faster than t-SNE, preserves global structure
- **Cons:** Many hyperparameters

#### 4. Autoencoders

- **Use Case:** Non-linear dimensionality reduction
  - **Pros:** Very flexible
  - **Cons:** Requires neural network training
- 

## Reinforcement Learning

### Key Concepts

- **Agent:** The learner/decision maker
- **Environment:** What the agent interacts with
- **State:** Current situation
- **Action:** What the agent can do
- **Reward:** Feedback from environment
- **Policy:** Strategy for choosing actions

### Common Algorithms

1. **Q-Learning:** Value-based method
  2. **Deep Q-Network (DQN):** Q-learning with neural networks
  3. **Policy Gradient:** Direct policy optimization
  4. **Actor-Critic:** Combines value and policy methods
  5. **PPO (Proximal Policy Optimization):** Stable policy gradient
- 

## Data Preprocessing Techniques

### 1. Handling Missing Data

- **Deletion:** Remove rows/columns with missing values
- **Imputation:**
  - Mean/Median/Mode imputation
  - Forward/Backward fill
  - KNN imputation
  - MICE (Multiple Imputation)

## 2. Feature Scaling

- **Standardization (Z-score):** Mean=0, SD=1

$$z = (x - \mu) / \sigma$$

- **Min-Max Normalization:** Scale to [0,1]

$$x_{\text{scaled}} = (x - \min) / (\max - \min)$$

- **Robust Scaling:** Using median and IQR

## 3. Encoding Categorical Variables

- **Label Encoding:** Ordinal categories
- **One-Hot Encoding:** Nominal categories
- **Target Encoding:** Based on target variable
- **Binary Encoding:** For high cardinality

## 4. Handling Imbalanced Data

- **Oversampling:** SMOTE, ADASYN
- **Undersampling:** Random, Tomek Links
- **Class Weights:** Adjust algorithm weights
- **Ensemble Methods:** BalancedRandomForest

## 5. Outlier Detection & Treatment

- **Statistical Methods:** Z-score, IQR
- **Isolation Forest**
- **Local Outlier Factor (LOF)**
- **Treatment:** Cap, transform, or remove

---

# Feature Engineering

## 1. Feature Creation

- **Polynomial Features:**  $x^2$ ,  $x^3$ ,  $x_1 \times x_2$
- **Interaction Features:** Combining features
- **Domain-Specific Features:** Based on expertise

## 2. Feature Transformation

- **Log Transformation:** For skewed data

- **Box-Cox Transformation:** Normalize distributions
- **Binning:** Continuous to categorical

### 3. Feature Selection

- **Filter Methods:**
  - Correlation analysis
  - Chi-square test
  - Mutual information
- **Wrapper Methods:**
  - Forward selection
  - Backward elimination
  - Recursive Feature Elimination (RFE)
- **Embedded Methods:**
  - Lasso (L1) regularization
  - Tree-based importance

### 4. Feature Extraction

- **PCA:** Linear combinations
  - **LDA:** Maximizes class separation
  - **Autoencoders:** Non-linear extraction
- 

## Model Evaluation Metrics

### Classification Metrics

#### 1. Accuracy

$$\text{Accuracy} = (\text{TP} + \text{TN}) / (\text{TP} + \text{TN} + \text{FP} + \text{FN})$$

#### 2. Precision

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

#### 3. Recall (Sensitivity)

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

#### 4. F1-Score

$$F1 = 2 \times (\text{Precision} \times \text{Recall}) / (\text{Precision} + \text{Recall})$$

## 5. ROC-AUC

- Area Under ROC Curve
- Trade-off between TPR and FPR

## 6. Confusion Matrix

- Visualizes TP, TN, FP, FN

## Regression Metrics

### 1. Mean Absolute Error (MAE)

$$MAE = (1/n) \times \sum |y_{\text{true}} - y_{\text{pred}}|$$

### 2. Mean Squared Error (MSE)

$$MSE = (1/n) \times \sum (y_{\text{true}} - y_{\text{pred}})^2$$

### 3. Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{MSE}$$

### 4. R<sup>2</sup> Score

$$R^2 = 1 - (SS_{\text{res}} / SS_{\text{tot}})$$

## Clustering Metrics

- **Silhouette Score:** Cohesion vs separation
- **Davies-Bouldin Index:** Cluster similarity
- **Calinski-Harabasz Index:** Ratio of dispersions

---

## Model Selection & Validation

### 1. Train-Test Split

- Typical split: 70-30 or 80-20
- Stratified split for imbalanced data



## 2. Cross-Validation

- **k-Fold**: Data split into k folds
- **Stratified k-Fold**: Maintains class distribution
- **Leave-One-Out (LOO)**:  $k = n$
- **Time Series Split**: For temporal data

## 3. Hyperparameter Tuning

- **Grid Search**: Exhaustive search
- **Random Search**: Random combinations
- **Bayesian Optimization**: Probabilistic model
- **Genetic Algorithms**: Evolutionary approach

## 4. Model Selection Criteria

- **Bias-Variance Trade-off**
  - **Occam's Razor**: Simpler is better
  - **Cross-validation scores**
  - **Business requirements**
- 

# Common Challenges & Solutions

## 1. Overfitting

**Symptoms**: High training accuracy, low test accuracy **Solutions**:

- Regularization (L1, L2)
- Dropout (neural networks)
- Early stopping
- More training data
- Simpler models
- Cross-validation

## 2. Underfitting

**Symptoms**: Low training and test accuracy **Solutions**:

- More complex models
- Feature engineering
- Reduce regularization

- Increase training time

### 3. Data Leakage

#### Prevention:

- Proper train-test split
- Avoid using future information
- Careful feature engineering

### 4. Curse of Dimensionality

#### Solutions:

- Feature selection
- Dimensionality reduction
- Regularization

### 5. Computational Constraints

#### Solutions:

- Feature selection
  - Model simplification
  - Distributed computing
  - Incremental learning
- 

## Deep Learning Overview

### Neural Network Basics

- **Perceptron:** Single neuron
- **Multi-Layer Perceptron (MLP):** Feedforward network
- **Activation Functions:** ReLU, Sigmoid, Tanh

### Common Architectures

#### 1. Convolutional Neural Networks (CNN)

- **Use Case:** Image processing
- **Key Components:** Conv layers, pooling, filters

#### 2. Recurrent Neural Networks (RNN)

- **Use Case:** Sequential data

- **Variants:** LSTM, GRU

### 3. Transformers

- **Use Case:** NLP, sequential data
- **Key Innovation:** Self-attention mechanism

### 4. Autoencoders

- **Use Case:** Dimensionality reduction, denoising
- **Types:** Vanilla, Variational (VAE)

### 5. Generative Adversarial Networks (GAN)

- **Use Case:** Data generation
- **Components:** Generator, Discriminator

## Deep Learning Best Practices

- **Batch Normalization:** Stabilize training
  - **Dropout:** Prevent overfitting
  - **Learning Rate Scheduling:** Adaptive learning
  - **Transfer Learning:** Use pre-trained models
  - **Data Augmentation:** Increase dataset size
- 

## Quick Decision Guide

### Choosing an Algorithm

#### For Classification:

- Small dataset → Naive Bayes, SVM
- Interpretability needed → Decision Tree, Logistic Regression
- High accuracy → Ensemble methods (Random Forest, XGBoost)
- Text data → Naive Bayes, SVM, Deep Learning

#### For Regression:

- Linear relationship → Linear Regression
- Non-linear → Random Forest, SVR, Neural Networks
- Interpretability → Linear Regression, Decision Trees

#### For Clustering:

- Known k → K-Means
- Unknown k → DBSCAN, Hierarchical
- Large dataset → Mini-batch K-Means

### For Dimensionality Reduction:

- Linear → PCA
- Visualization → t-SNE, UMAP
- Feature selection → Lasso, Tree importance

### Performance Optimization Checklist

1. ✓ Data quality check
  2. ✓ Appropriate preprocessing
  3. ✓ Feature engineering
  4. ✓ Algorithm selection
  5. ✓ Hyperparameter tuning
  6. ✓ Cross-validation
  7. ✓ Ensemble methods
  8. ✓ Error analysis
  9. ✓ Business metric alignment
  10. ✓ Model monitoring
- 

### Resources for Implementation

#### Python Libraries

- **General ML:** scikit-learn, XGBoost, LightGBM
- **Deep Learning:** TensorFlow, PyTorch, Keras
- **Data Processing:** pandas, NumPy
- **Visualization:** matplotlib, seaborn, plotly
- **AutoML:** auto-sklearn, H2O, TPOT

### Key Formulas Reference

- **Bias:**  $E[\hat{f}(x)] - f(x)$
- **Variance:**  $E[(\hat{f}(x) - E[\hat{f}(x)])^2]$
- **Information Gain:**  $H(\text{parent}) - \sum(\text{weighted } H(\text{children}))$
- **Gini Index:**  $1 - \sum(p_i)^2$

- **Entropy:**  $-\sum(p_i \times \log_2(p_i))$